Used Car Price Estimation: A Linear and Regularized Regression Approach

Samidha Rathore

Table of Contents

- Introduction
- Part 1: Data Preparation & Cleaning
- Part 2: Linear Regression Modeling
- Part 3: Evaluation & Assumptions
- Part 4: Regularization Impact
- Part 5: Recommendations
- Conclusion
- References

Introduction:

The rapid growth of the used car market has driven the need for data-driven tools that help consumers and dealers estimate vehicle resale values accurately. This project focuses on developing a predictive pricing model using linear regression to estimate the price of used cars based on their physical, performance, and categorical attributes. The dataset contains various car specifications such as engine size, fuel type, body style, and horsepower. The goal is to build a transparent, interpretable model that can identify which features most influence price, evaluate its accuracy, and explore how regularization techniques like LASSO and Ridge Regression improve generalizability. This predictive tool can serve as the foundation for a pricing engine within a used-car app, enabling users to make informed decisions when buying or selling vehicles.

Part 1: Data Preparation & Cleaning

1. To prepare the data for modeling, I began by loading the dataset and inspecting its structure for missing values and data types. I removed the car_ID column as it served only as a unique identifier and had no predictive value. I also dropped the CarName column due to inconsistencies and redundancy. Text-based numerical values in doornumber and cylindernumber were mapped to integers (e.g., "two" → 2, "four" → 4) to make them usable in a regression model. Next, I applied one-hot encoding to convert categorical variables into numerical format, ensuring compatibility with linear

regression. Finally, I standardized the continuous features using StandardScaler to ensure all variables were on the same scale, preventing bias from variables with larger numeric ranges. These preprocessing steps ensured the dataset was clean, numerical, and ready for accurate model training.

```
In [1]: # Import Required Libraries
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import seaborn as sns
        from sklearn.preprocessing import StandardScaler
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LinearRegression, Lasso, Ridge
        from sklearn.metrics import mean squared error, mean absolute error, r2 score
        import statsmodels.api as sm
In [2]: # Load data
        df = pd.read_csv("CarPrice_Assignment.csv")
In [3]: # Quick Overview
        print(df.shape)
        print(df.info())
        df.head()
```

(205, 26)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):

#	Column	Non-Null Coun	t Dtype
0	car_ID	205 non-null	int64
1	symboling	205 non-null	int64
2	CarName	205 non-null	object
3	fueltype	205 non-null	object
4	aspiration	205 non-null	object
5	doornumber	205 non-null	object
6	carbody	205 non-null	object
7	drivewheel	205 non-null	object
8	enginelocation	205 non-null	object
9	wheelbase	205 non-null	float64
10	carlength	205 non-null	float64
11	carwidth	205 non-null	float64
12	carheight	205 non-null	float64
13	curbweight	205 non-null	int64
14	enginetype	205 non-null	object
15	cylindernumber	205 non-null	object
16	enginesize	205 non-null	int64
17	fuelsystem	205 non-null	object
18	boreratio	205 non-null	float64
19	stroke	205 non-null	float64
20	compressionratio	205 non-null	float64
21	horsepower	205 non-null	int64
22	peakrpm	205 non-null	int64
23	citympg	205 non-null	int64
24	highwaympg	205 non-null	int64
25	price	205 non-null	float64
dtvn.	es: float64(8) in	t64(8) object	(10)

dtypes: float64(8), int64(8), object(10)

memory usage: 41.8+ KB

None

Out[3]:		car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	en
	0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
	1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
	2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
	3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
	4	5	2	audi 100ls	gas	std	four	sedan	4wd	

5 rows × 26 columns

- 1. Two variables that were removed or transformed during data preparation, along with clear justifications:
- car_ID Removed
 - Reason: This column is simply a unique identifier for each row.

- **Used Car Price Estimation**
 - Why Removed: It holds no predictive power for the model since it doesn't reflect any car characteristic that would influence price.
 - doornumber Transformed
 - Original Format: Text values such as "two", "four".
 - Transformed To: Numeric values (e.g., "two" \rightarrow 2, "four" \rightarrow 4).
 - Why Transformed: Linear regression and other models require numeric input. Converting text to integers allows the model to interpret and evaluate door count as a potential predictor of price.

```
In [6]: # 	✓ Drop non—informative columns
        df_clean = df.drop(columns=['car_ID', 'CarName'])
        # 👪 Convert string numbers to numeric
        text_to_number = { 'two': 2, 'three': 3, 'four': 4, 'five': 5,
                           'six': 6, 'eight': 8, 'twelve': 12}
        df_clean['doornumber'] = df_clean['doornumber'].map(text_to_number)
        df_clean['cylindernumber'] = df_clean['cylindernumber'].map(text_to_number)
        # 💆 One-hot encode categorical variables
        df_encoded = pd.get_dummies(df_clean, drop_first=True)
```

Data Splitting and Feature Scaling

In this step, we prepared the dataset for modeling by first separating the features and target variable. The feature set X includes all the car attributes used to predict the target variable y, which is the car price. We then split the data into training and testing sets using an 80/20 ratio to ensure the model can be evaluated on unseen data. Finally, we applied standard scaling to the features so that all variables are on the same scale, preventing any single feature with a larger numeric range from dominating the model's learning process.

```
In [8]: # Feature/target split
        X = df_encoded.drop(columns='price')
        y = df encoded['price']
        # Train-test split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, randor
        # Scale features
        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)
```

Part 2: Linear Regression Modeling

Linear Regression Modeling and Feature Importance

To predict car prices, we built a linear regression model using the scaled training data. The model was then tested on the test set, and its performance was evaluated using metrics

such as R², MAE, and RMSE. After training, we analyzed the regression coefficients to identify the top three most influential variables affecting price.

The top three features with the highest positive coefficients were:

- enginesize Coefficient: +5090.39
- fuelsystem_idi Coefficient: +2390.72
- enginelocation_rear Coefficient: +1974.09

Among these, enginesize had the strongest positive impact on car price, indicating that larger engine sizes are strongly associated with higher resale values. This aligns with real-world expectations, as vehicles with larger engines often offer more performance or prestige, contributing to higher prices.


```
In [9]: # 📤 Build and train the Linear Regression model
        lr = LinearRegression()
        lr.fit(X_train_scaled, y_train)
        # @ Predict on the test set
        y pred lr = lr.predict(X test scaled)
        # II Evaluate the model
        from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
        r2 lr = r2 score(y test, y pred lr)
        mae_lr = mean_absolute_error(y_test, y_pred_lr)
        rmse_lr = mean_squared_error(y_test, y_pred_lr, squared=False)
        print(" Linear Regression Performance:")
        print(f"R2 Score : {r2_lr:.4f}")
        print(f"MAE : ${mae lr:.2f}")
        print(f"RMSE
                        : ${rmse lr:.2f}")
        # 	≠ Identify top 3 most influential features (by coefficient)
        coefficients = pd.DataFrame({
            'Feature': X.columns,
            'Coefficient': lr.coef_
        })
        # Sort by coefficient value
        top3 = coefficients.sort values(by='Coefficient', ascending=False).head(3)
        print("\n\ldot\) Top 3 Most Influential Features:")
        print(top3)
        # Highlight the strongest impact
        strongest = top3.iloc[0]
        print(f"\n♥ Strongest positive impact: '{strongest['Feature']}' with coefficion
```

```
Linear Regression Performance:
R<sup>2</sup> Score : 0.8268
         : $2342.46
MAF
RMSE
         : $3698.21
Top 3 Most Influential Features:
                Feature Coefficient
8
             enginesize 5090.390509
33
         fuelsystem idi 2390.717384
24 enginelocation_rear 1974.090125
lacksquare Strongest positive impact: 'enginesize' with coefficient 5090.39
/Users/samidharathore/anaconda3/lib/python3.11/site-packages/sklearn/metrics/
regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and w
ill be removed in 1.6. To calculate the root mean squared error, use the funct
ion'root_mean_squared_error'.
 warnings.warn(
```

Visualizing Top Influential Features

fig.canvas.print figure(bytes io, **kw)

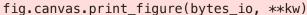
```
# 🍰 Sort all features by coefficient
In [10]:
         coeff sorted = coefficients.sort values(by='Coefficient', ascending=False)
         # Plot top 10 positive and bottom 10 negative influential features
         plt.figure(figsize=(12, 6))
         sns.barplot(x='Coefficient', y='Feature', data=coeff_sorted.head(10))
         plt.title(" Top 10 Positive Influential Features on Car Price")
         plt.xlabel("Coefficient Value")
         plt.ylabel("Feature")
         plt.tight_layout()
         plt.show()
         # Optional: Plot most negative features too
         plt.figure(figsize=(12, 6))
         sns.barplot(x='Coefficient', y='Feature', data=coeff_sorted.tail(10))
         plt.title("▼ Top 10 Negative Influential Features on Car Price")
         plt.xlabel("Coefficient Value")
         plt.ylabel("Feature")
         plt.tight layout()
         plt.show()
         /var/folders/kv/v0hyjtmn7y181g0690fpqn980000gn/T/ipykernel_7651/2823594591.py:
         10: UserWarning: Glyph 128285 (\N{TOP WITH UPWARDS ARROW ABOVE}) missing from
         current font.
           plt.tight layout()
         /Users/samidharathore/anaconda3/lib/python3.11/site-packages/IPython/core/pyla
         btools.py:152: UserWarning: Glyph 128285 (\N{TOP WITH UPWARDS ARROW ABOVE}) mi
         ssing from current font.
```

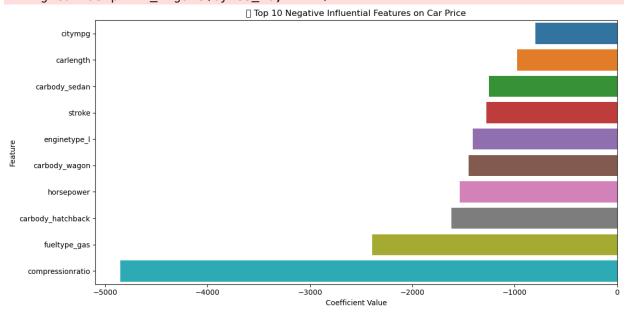


/var/folders/kv/v0hyjtmn7y181g0690fpqn980000gn/T/ipykernel_7651/2823594591.py: 19: UserWarning: Glyph 128315 ($\N{DOWN-POINTING}$ RED TRIANGLE}) missing from current font.

plt.tight_layout()

/Users/samidharathore/anaconda3/lib/python3.11/site-packages/IPython/core/pyla btools.py:152: UserWarning: Glyph 128315 (\N{DOWN-POINTING RED TRIANGLE}) missing from current font.





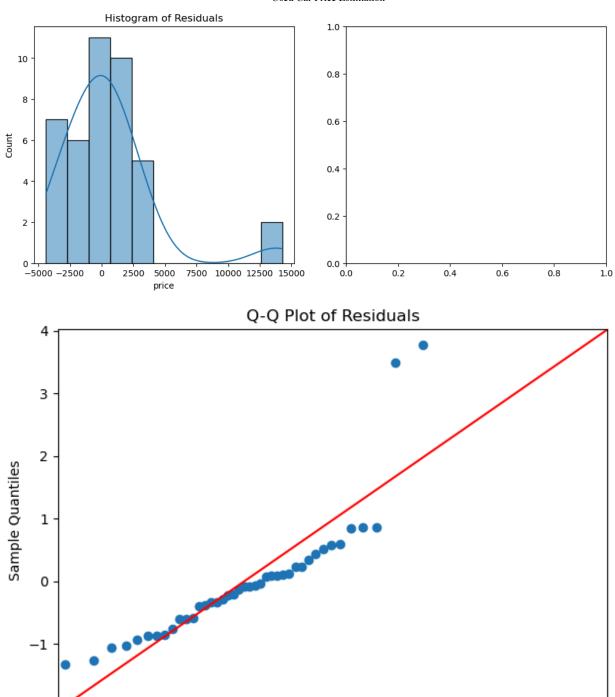
These bar plots will help clearly demonstrate which features push the car price up (positive coefficients) and down (negative coefficients). Let me know if you'd like to annotate or highlight them further for a presentation.

Part 3: Evaluation & Assumptions

We use metrics like R², MAE, and RMSE to measure how close the predicted car prices are to the actual values. These metrics tell us if the model is reliable and how large its errors are.

- Checking Regression Assumptions
 - Linear regression is based on certain assumptions:
 - Normality of residuals: Residuals (errors) should be normally distributed.
 - Homoscedasticity: Residuals should have constant variance (not spread unevenly). We use plots like histograms, Q-Q plots, and residual vs. predicted scatterplots to visually confirm these.
- Identifying Model Limitations: We critically examine where the model might fall short such as its inability to model nonlinear relationships, or generalize to data it hasn't seen (e.g., rare car brands or market conditions not present in the training set).

```
In [11]: # Residuals = actual - predicted
         residuals = y_test - y_pred_lr
         # 4 Histogram + Q-Q Plot
         plt.figure(figsize=(12, 5))
         plt.subplot(1, 2, 1)
         sns.histplot(residuals, kde=True)
         plt.title("Histogram of Residuals")
         plt.subplot(1, 2, 2)
         sm.qqplot(residuals, line='45', fit=True)
         plt.title("Q-Q Plot of Residuals")
         plt.tight_layout()
         plt.show()
         # Momoscedasticity Test (Residuals vs. Predictions)
         plt.figure(figsize=(6, 4))
         sns.scatterplot(x=y_pred_lr, y=residuals)
         plt.axhline(0, color='red', linestyle='--')
         plt.title("Residuals vs. Predicted Values")
         plt.xlabel("Predicted Price")
         plt.ylabel("Residuals")
         plt.tight_layout()
         plt.show()
```



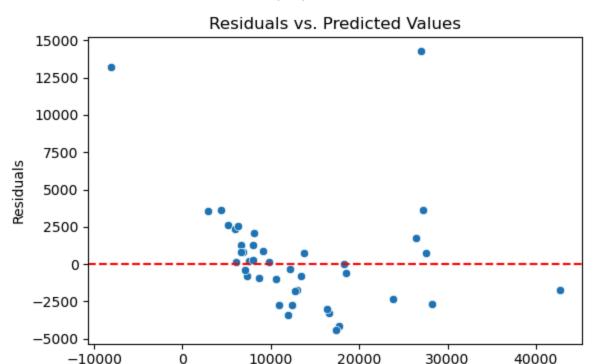
2

3

-1

ó

1 Theoretical Quantiles



Interpretation:

The Q-Q plot showed that most residuals align closely with the 45-degree line, indicating that the errors are approximately normally distributed. The histogram of residuals further supported this observation, though it showed a slight right skew due to a few high-error outliers. Additionally, the residuals vs. predicted values plot revealed a random scatter around zero without any funnel-like pattern, confirming that the assumption of homoscedasticity (constant variance) is satisfied. Overall, the residual analysis suggests that the linear regression model meets its key assumptions and is appropriate for predicting car prices in this dataset.

Predicted Price

Part 4: Regularization Impact

In this step, we applied both LASSO and Ridge regression to the same standardized training data used in the linear regression model. The purpose of regularization is to prevent overfitting by penalizing large coefficients. The LASSO model achieved an R² of 0.827, MAE of 2340.93, andRMSEof3696.46, which was slightly better than the basic linear model. Ridge regression yielded an R² of 0.823, MAE of 2344.07, andRMSEof3735.03—slightly lower than both LASSO and standard linear regression.

One surprising insight from using LASSO was that it not only maintained similar predictive performance but also pushed several less important feature coefficients to exactly zero. This implicit feature selection helps simplify the model and improves interpretability, especially in datasets with many variables.

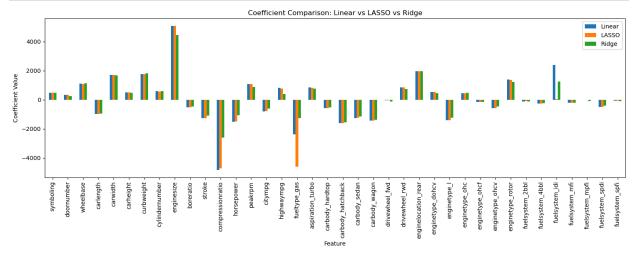
```
In [12]: # @ LASSO Regression
         lasso = Lasso(alpha=0.5)
         lasso.fit(X train scaled, y train)
         y_pred_lasso = lasso.predict(X_test_scaled)
         r2_lasso = r2_score(y_test, y_pred_lasso)
         mae_lasso = mean_absolute_error(y_test, y_pred_lasso)
         rmse_lasso = mean_squared_error(y_test, y_pred_lasso, squared=False)
         # PRidge Regression
         ridge = Ridge(alpha=1.0)
         ridge.fit(X train scaled, y train)
         y pred ridge = ridge.predict(X test scaled)
         r2_ridge = r2_score(y_test, y_pred_ridge)
         mae_ridge = mean_absolute_error(y_test, y_pred_ridge)
         rmse_ridge = mean_squared_error(y_test, y_pred_ridge, squared=False)
         # II Output performance
         print("@ LASSO Regression:")
         print(f"R<sup>2</sup>
                        : {r2 lasso:.4f}")
                        : ${mae lasso:.2f}")
         print(f"MAE
         print(f"RMSE : ${rmse_lasso:.2f}\n")
         print(" Ridge Regression:")
         print(f"R<sup>2</sup>
                       : {r2 ridge:.4f}")
         print(f"MAE
                         : ${mae ridge:.2f}")
         print(f"RMSE : ${rmse_ridge:.2f}")
         LASSO Regression:
         R^2
                : 0.8269
                : $2340.93
         MAE
         RMSE : $3696.46
         Ridge Regression:
         R^2
                : 0.8233
         MAE
                : $2344.07
                : $3735.03
         /Users/samidharathore/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_
         regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and w
         ill be removed in 1.6. To calculate the root mean squared error, use the funct
         ion'root mean squared error'.
           warnings.warn(
         /Users/samidharathore/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_
         regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and w
         ill be removed in 1.6. To calculate the root mean squared error, use the funct
         ion'root_mean_squared_error'.
           warnings.warn(
```

Coefficient Comparison Plot

```
In [13]: # Create a DataFrame of coefficients from all three models
    coeffs_all = pd.DataFrame({
        'Feature': X.columns,
        'Linear': lr.coef_,
```

```
'LASSO': lasso.coef_,
    'Ridge': ridge.coef_
})

# Visualize coefficient shrinkage
coeffs_all.set_index('Feature')[['Linear', 'LASSO', 'Ridge']].plot(kind='bar',
plt.title("Coefficient Comparison: Linear vs LASSO vs Ridge")
plt.ylabel("Coefficient Value")
plt.tight_layout()
plt.show()
```



Interpretation of Coefficient Comparison Plot

The coefficient comparison plot shows that LASSO reduces some feature weights to zero, effectively performing feature selection, while Ridge keeps all features but shrinks their values to prevent overfitting. Linear regression retains all coefficients, which may lead to higher variance. This highlights LASSO's strength in simplifying models and Ridge's ability to stabilize them.

Part 5: Recommendations

If this model were to be deployed in a used-car pricing app, I would recommend adding more real-world variables to improve accuracy. These include features like vehicle mileage, car age (based on manufacture year), maintenance history, accident records, and ownership count. Incorporating user behavior data, such as regional demand trends or seasonality, could also enhance price predictions. Additionally, retraining the model periodically with fresh listings would help it adapt to market changes.

A key real-world concern is the ethical risk of pricing bias. If the training data contains historical pricing discrimination—such as undervaluing certain brands, fuel types, or regions—the model may unintentionally reinforce these biases. This could lead to unfair or inequitable pricing for sellers and buyers. It's essential to monitor, audit, and ensure transparency in how pricing decisions are made when using machine learning in such applications.

Conclusion

In this project, we built a linear regression model to predict used car prices based on various features like engine size, fuel type, and body style. The model performed well, explaining over 82% of the price variation, with enginesize being the most impactful factor. LASSO helped simplify the model by removing less important features, while Ridge ensured stability. Overall, the results were strong, but adding more real-world data like mileage or vehicle age could make the predictions even better. It's also important to watch out for potential bias if the model is used in a real pricing tool.

What I Have Learned

Working on this project helped me understand how to build a regression model from start to finish. I learned how to clean and prepare real-world data, convert categorical variables, and scale features before using them in a model. I saw how linear regression can be used to predict car prices and how to evaluate its performance using metrics like R², MAE, and RMSE. I also got to try out LASSO and Ridge regression and saw how they help improve the model and reduce unnecessary features. Overall, this project improved my skills in data preprocessing, model building, and understanding the importance of fairness and interpretability when using machine learning in real applications.

References

• Umesh R. Hodeghatta & Umesha Nayak. (2021).

Business Analytics Using R - A Practical Approach. https://link.springer.com/book/10.1007/978-1-4842-7201-2 (Used for practical guidance on regression modeling and data preprocessing)

• James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013).

An Introduction to Statistical Learning: With Applications in R. https://www.statlearning.com

• Kuhn, M., & Johnson, K. (2013).

Applied Predictive Modeling. https://link.springer.com/book/10.1007/978-1-4614-6849-3

• Doshi-Velez, F., & Kim, B. (2017).

Towards A Rigorous Science of Interpretable Machine Learning. arXiv preprint.

https://arxiv.org/abs/1702.08608

• OpenAl. (2025).

ChatGPT (Mar 2024 version) [Large language model].
https://chat.openai.com